

Inside Visual C++'s Parallel Algorithms

What happens when you throw `std::execution::par...`

<https://github.com/BillyONeal/InsideParallelAlgorithms>

Billy O'Neal

Sr. SDE / Standard Library Maintainer @ Microsoft

bion@microsoft.com @MalwareMinigun

Hello Parallel World – Not This Talk



- See blog post
<https://blogs.msdn.microsoft.com/vcblog/2018/09/11/using-c17-parallel-algorithms-for-better-performance/>
- `std::sort(a, b) -> std::sort(std::execution::par, a, b)`
- 75ms -> 20ms (1000000 doubles on 7980XE)

So you want to parallelize an algorithm....

- Partitioning (not `std::partition!`)
- Schedule on compute resources
- Merge results

Let's look at accumulate...

X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++
X = X + *|t++



⋮

Partitioning...

```
X = X + *|t++  
X = X + *|t++  
X = X + *|t++
```

```
X = X + *|t++  
X = X + *|t++  
X = X + *|t++
```

```
X = X + *|t++  
X = X + *|t++  
X = X + *|t++
```

```
X = X + *|t++  
X = X + *|t++
```

Processor

Processor

...

Partitioning...

$X = X + *|t++$
 $X = X + *|t++$
 $X = X + *|t++$

$X = X + *|t++$
 $X = X + *|t++$
 $X = X + *|t++$

$X = X + *|t++$
 $X = X + *|t++$
 $X = X + *|t++$

$X = X + *|t++$
 $X = X + *|t++$

Data race on X, it!

Processor

Processor

...

Partitioning...

```
lx = lt  
X = X + *lx++  
X = X + *lx++  
X = X + *lx++
```

```
la = next(lt, 3)  
A = *la++ + *la++  
A = A + *la++
```

```
lb = next(lt, 6)  
B = *lb++ + *lb++  
B = B + *lb++
```

```
lc = next(lt, 9)  
C = *lc++ + *lc++
```

Processor

Processor

...

Merge Results...

```
lx = lt  
X = X + *lx++  
X = X + *lx++  
X = X + *lx++
```

```
la = next(lt, 3)  
A = *la++ + *la++  
A = A + *la++
```

```
lb = next(lt, 6)  
B = *lb++ + *lb++  
B = B + *lb++
```

```
lc = next(lt, 9)  
C = *lc++ + *lc++
```

```
(wait)  
X = X + A  
X = X + B  
X = X + C
```

Processor

Processor

⋮

```
lx = lt  
X = X + *lx++  
X = X + *lx++  
X = X + *lx++
```

```
la = next(lt, 3)  
A = *la++ + *la++  
A = A + *la++
```

```
lb = next(lt, 6)  
B = *lb++ + *lb++  
B = B + *lb++
```

```
lc = next(lt, 9)  
C = *lc++ + *lc++
```

```
(wait)  
X = X + A  
X = X + C  
X = X + B
```

Processor

Processor

⋮

Let's look at ~~accumulate~~ reduce...

What happened to scheduling?

```
lx = lt  
X = X + *lx++  
X = X + *lx++  
X = X + *lx++
```

```
la = next(lt)  
A = *la++  
A = A + 1
```

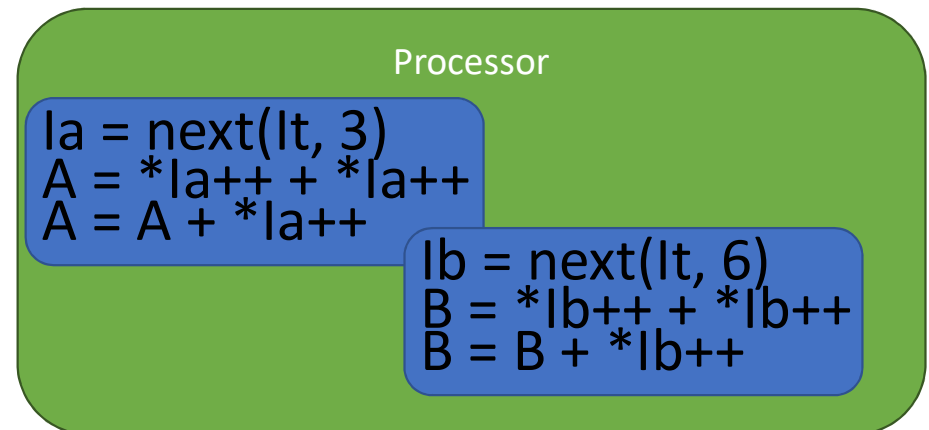
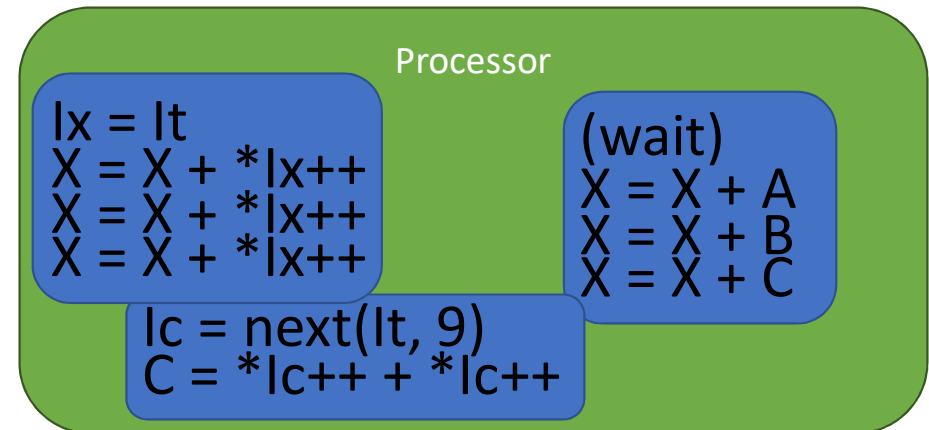
```
lb = next(la)  
B = *lb++  
B = B + 1
```

```
lc = next(lb)  
C = *lc++
```

```
(wait)  
X = X + A  
X = X + B  
X = X + C
```

Scheduling Magic

⋮



...

Scheduling, really – Windows' Thread Pool

- [CreateThreadpoolWork](#)
- [SubmitThreadpoolWork](#)
- [WaitForThreadpoolWorkCallbacks](#)
- [CloseThreadpoolWork](#)
- [Pedro Teixeira's talk](#) (Pedro works on the kernel) discusses threadpool internals

Questions on the mental model of reduce?

Demo – Debugging into `std::reduce`

Benchmark benchmark benchmark!

- The hardware and the input you care about are important
- Parallel algorithms generally do more work, even for `for_each(random-access)`, to acquire threads, wait for background threads to complete, etc.
- Following numbers are from this ThinkPad X1 Carbon, i7-8650U 4c8t, 2133MHz DDR3L

Benchmark benchmark benchmark - Debug

- .\DemoReduce.exe 1000000000 (762mb) – Parallel 4.9 times faster
- .\DemoReduce.exe 100000000 (76mb) – Parallel 4.6 times faster
- .\DemoReduce.exe 1000000 (7.6mb) – Parallel 3.8 times faster
(fits in cache on this chip below here)
- .\DemoReduce.exe 100000 (.76mb) – Parallel 1.5 times faster
- .\DemoReduce.exe 1000 – Parallel 13 times slower

Benchmark benchmark benchmark - Release

- .\DemoReduce.exe 1000000000 (762mb) – Parallel 1.3 times faster
- .\DemoReduce.exe 100000000 (76mb) – Parallel 1.3 times faster
- .\DemoReduce.exe 10000000 (7.6mb) – Parallel 1.6 times faster
(fits in cache on this chip below here)
- .\DemoReduce.exe 100000 (.76mb) – Parallel 2 times slower
- .\DemoReduce.exe 1000 – Parallel 96 times slower

Demo - Why isn't it any faster?

Hardware matters!

- Next from the 7980XE 18c36t; 3200MHz DDR4

Benchmark benchmark benchmark - Debug

- .\DemoReduce.exe 1000000000 (762mb) – Parallel 16.6 times faster
- .\DemoReduce.exe 100000000 (76mb) – Parallel 8.8 times faster
- (fits in cache on this chip below here)
- .\DemoReduce.exe 1000000 (7.6mb) – Parallel 3.8 times faster
- .\DemoReduce.exe 100000 (.76mb) – Parallel 2.8 times faster
- .\DemoReduce.exe 1000 – Parallel 15 times slower

Benchmark benchmark benchmark - Release

- .\DemoReduce.exe 1000000000 (762mb) – Parallel 4.7 times faster
- .\DemoReduce.exe 100000000 (76mb) – Parallel 5.3 times faster
- (fits in cache on this chip below here)
- .\DemoReduce.exe 1000000 (7.6mb) – Parallel 2.1 times faster
- .\DemoReduce.exe 100000 (.76mb) – Parallel 3.9 times slower
- .\DemoReduce.exe 1000 – Parallel 157 times slower

Questions about Benchmarks?

A more “interesting” algorithm –
stable_sort

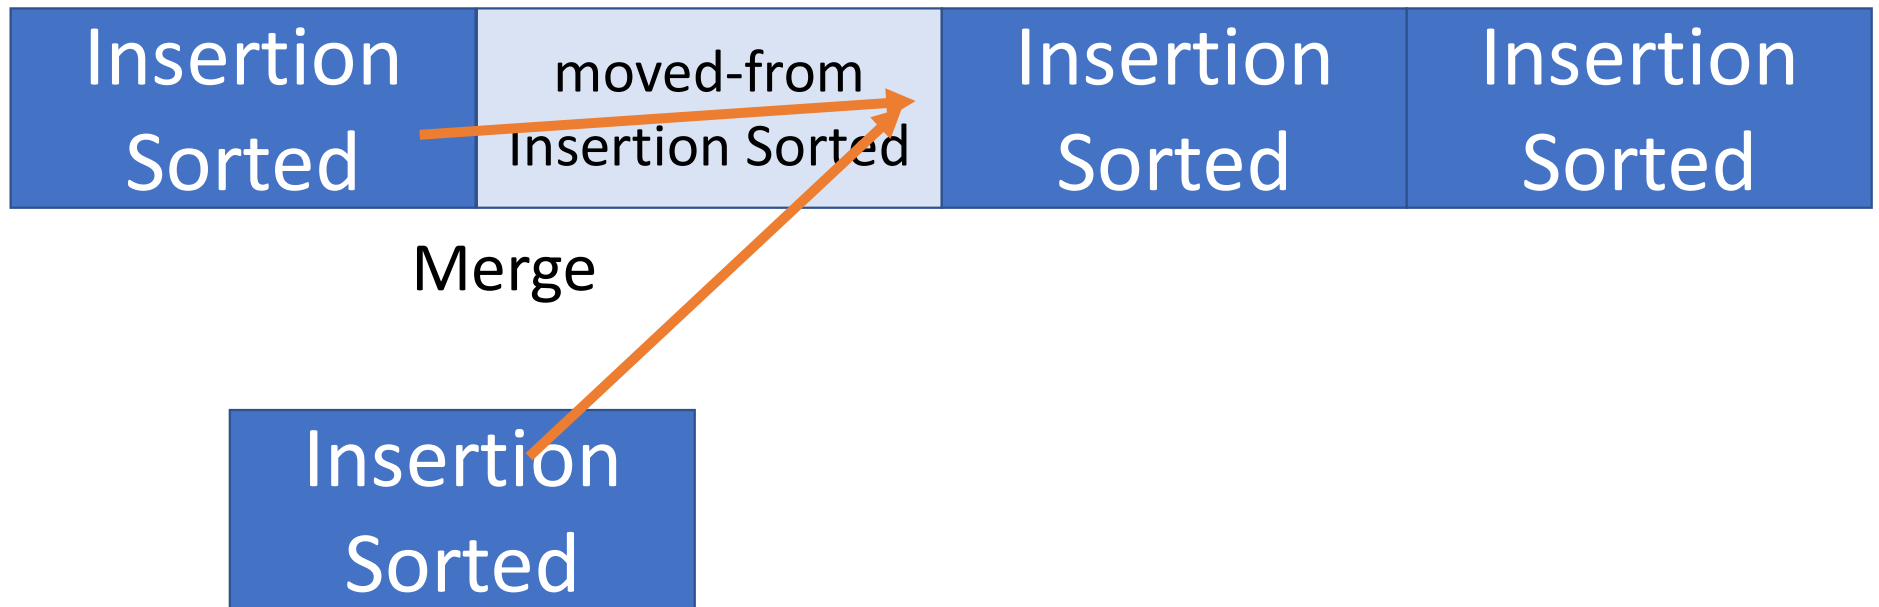
stable_sort

Unsorted

Insertion sort element chunks

Insertion Sorted	Insertion Sorted	Insertion Sorted	Insertion Sorted
---------------------	---------------------	---------------------	---------------------

inplace_merge elementwise



inplace_merge elementwise

Merged	Insertion Sorted	Insertion Sorted
--------	---------------------	---------------------

inplace_merge elementwise



Worst case space consumption, $n/2$

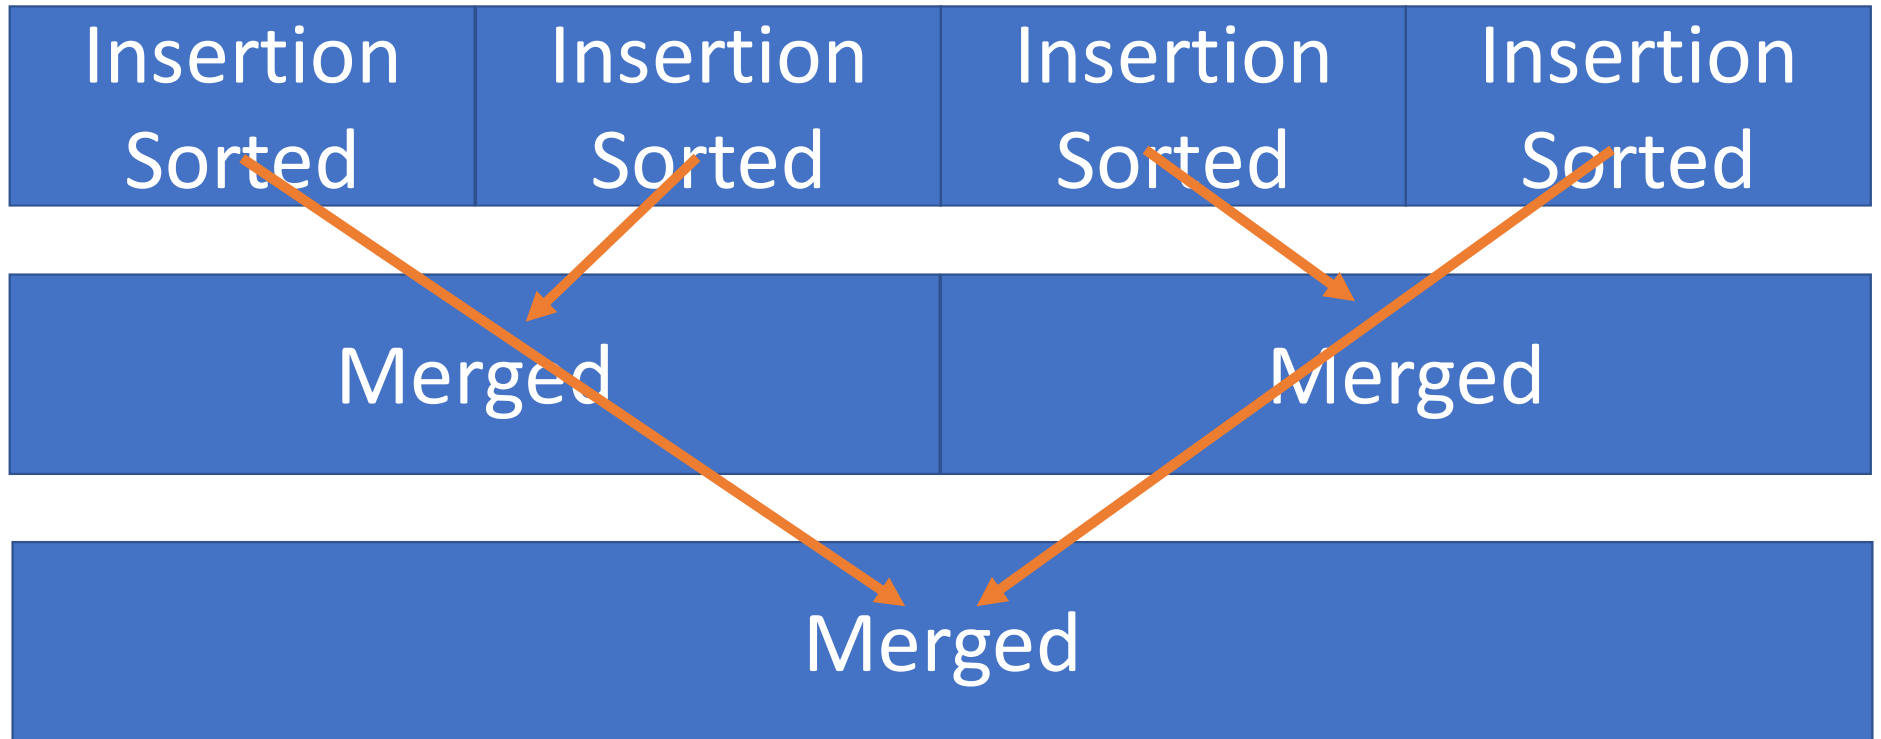
Sorted

Merged

Insertion Sorted	Insertion Sorted	Insertion Sorted	Insertion Sorted
---------------------	---------------------	---------------------	---------------------

Merged	Merged
--------	--------

Merged



Get temporary space




Unsorted



Temp Space

Partition input + temp space together

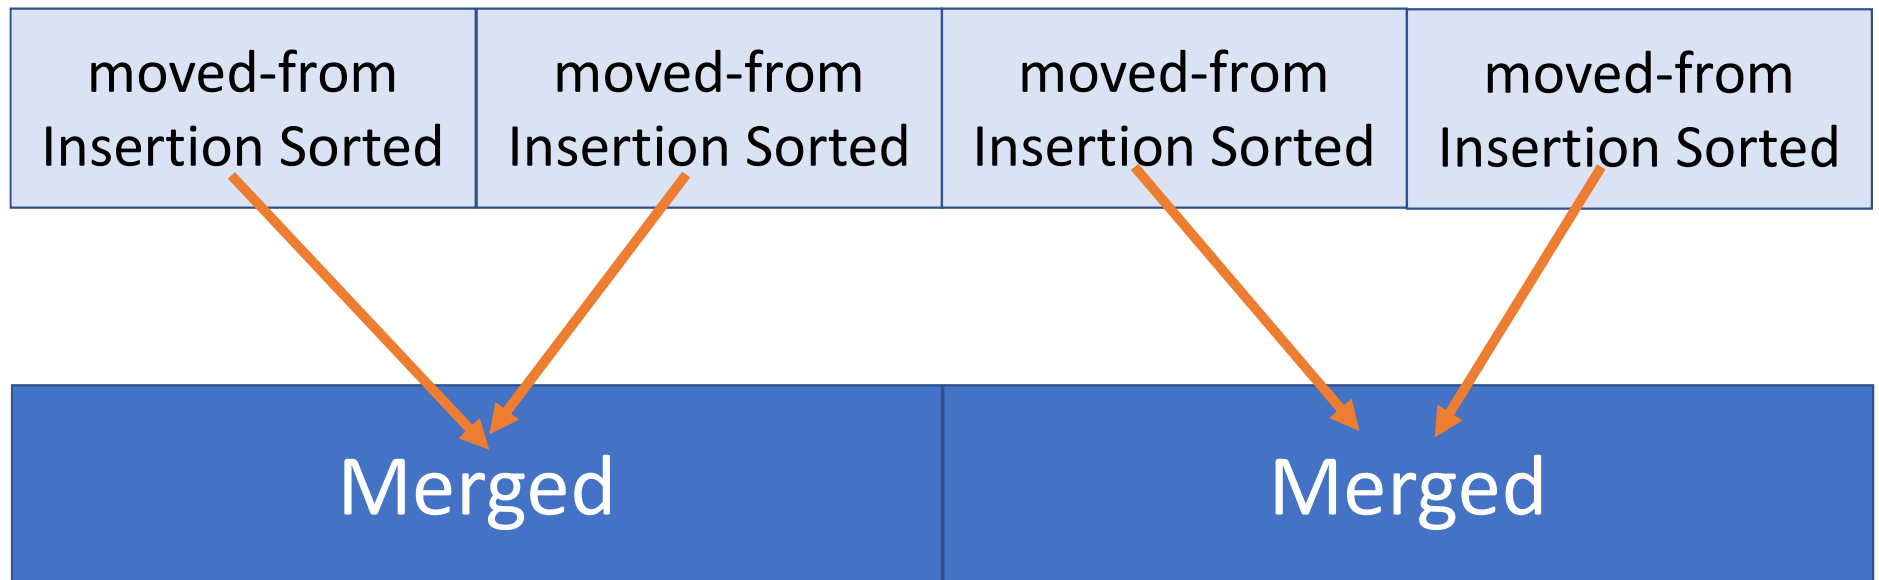


Insertion Sort – 4 Partitions


Insertion Sorted	Insertion Sorted	Insertion Sorted	Insertion Sorted
---------------------	---------------------	---------------------	---------------------

Temp Space

merge – 2 partitions

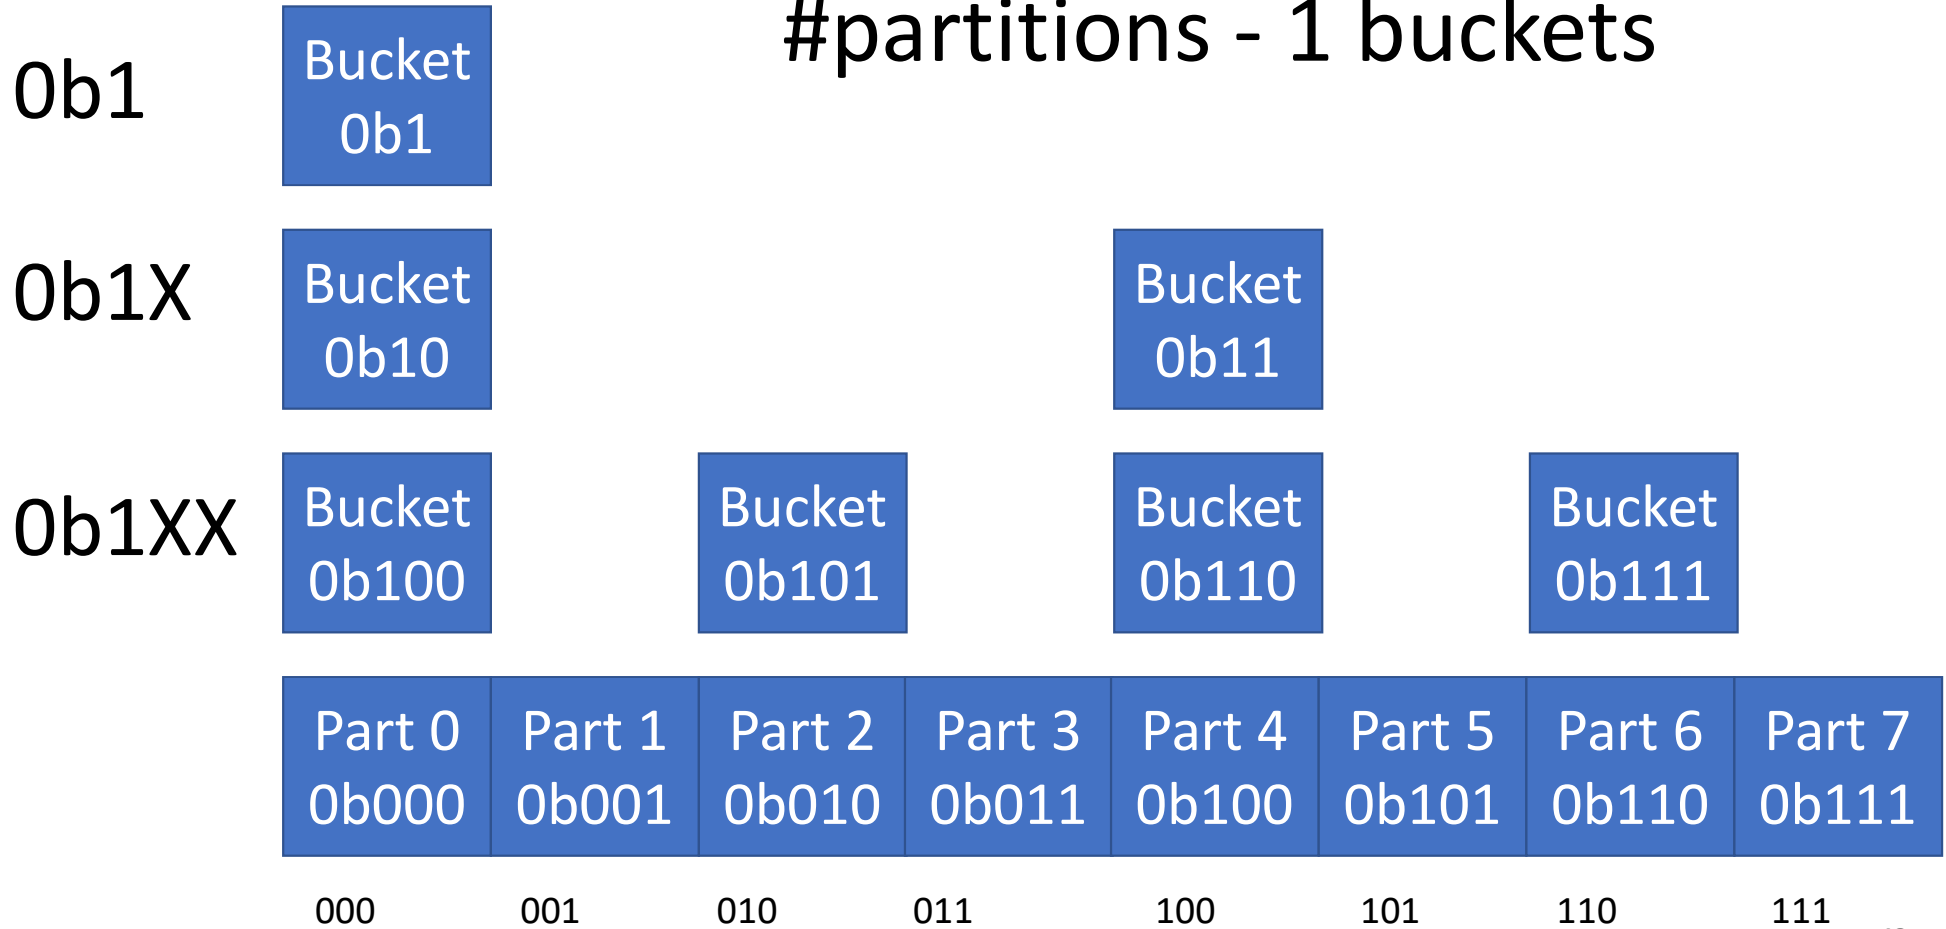


merge – 1 partition

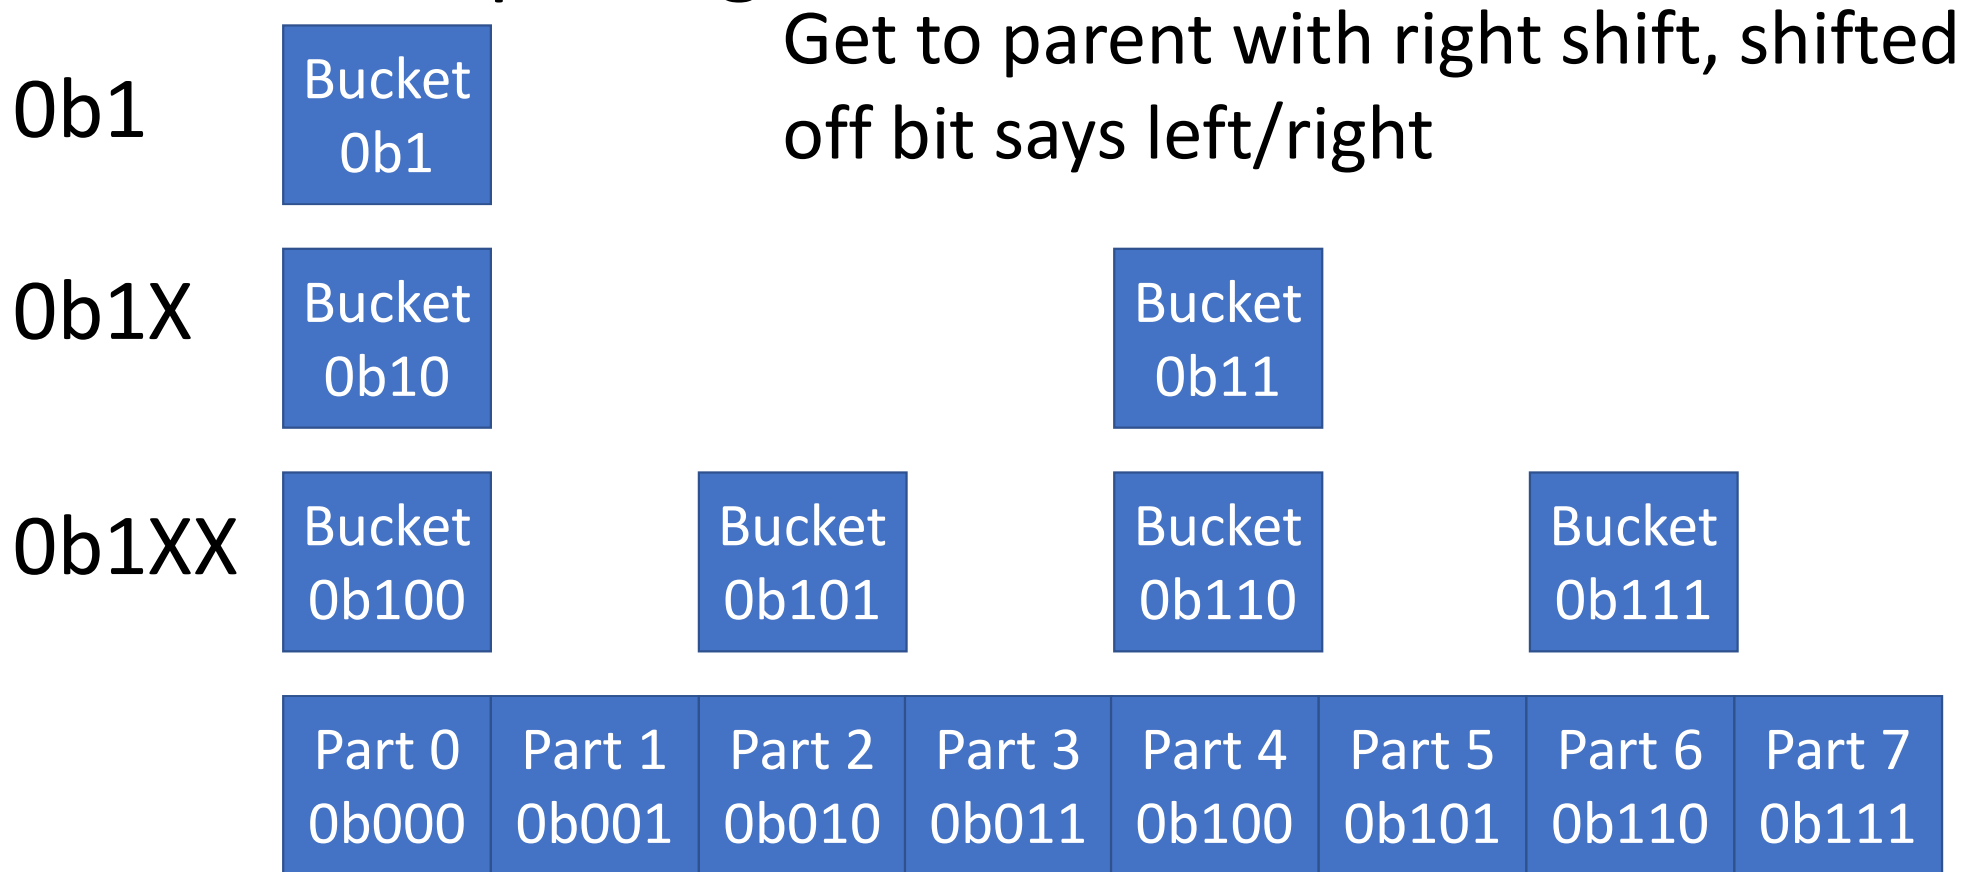


Bottom-up merge tree

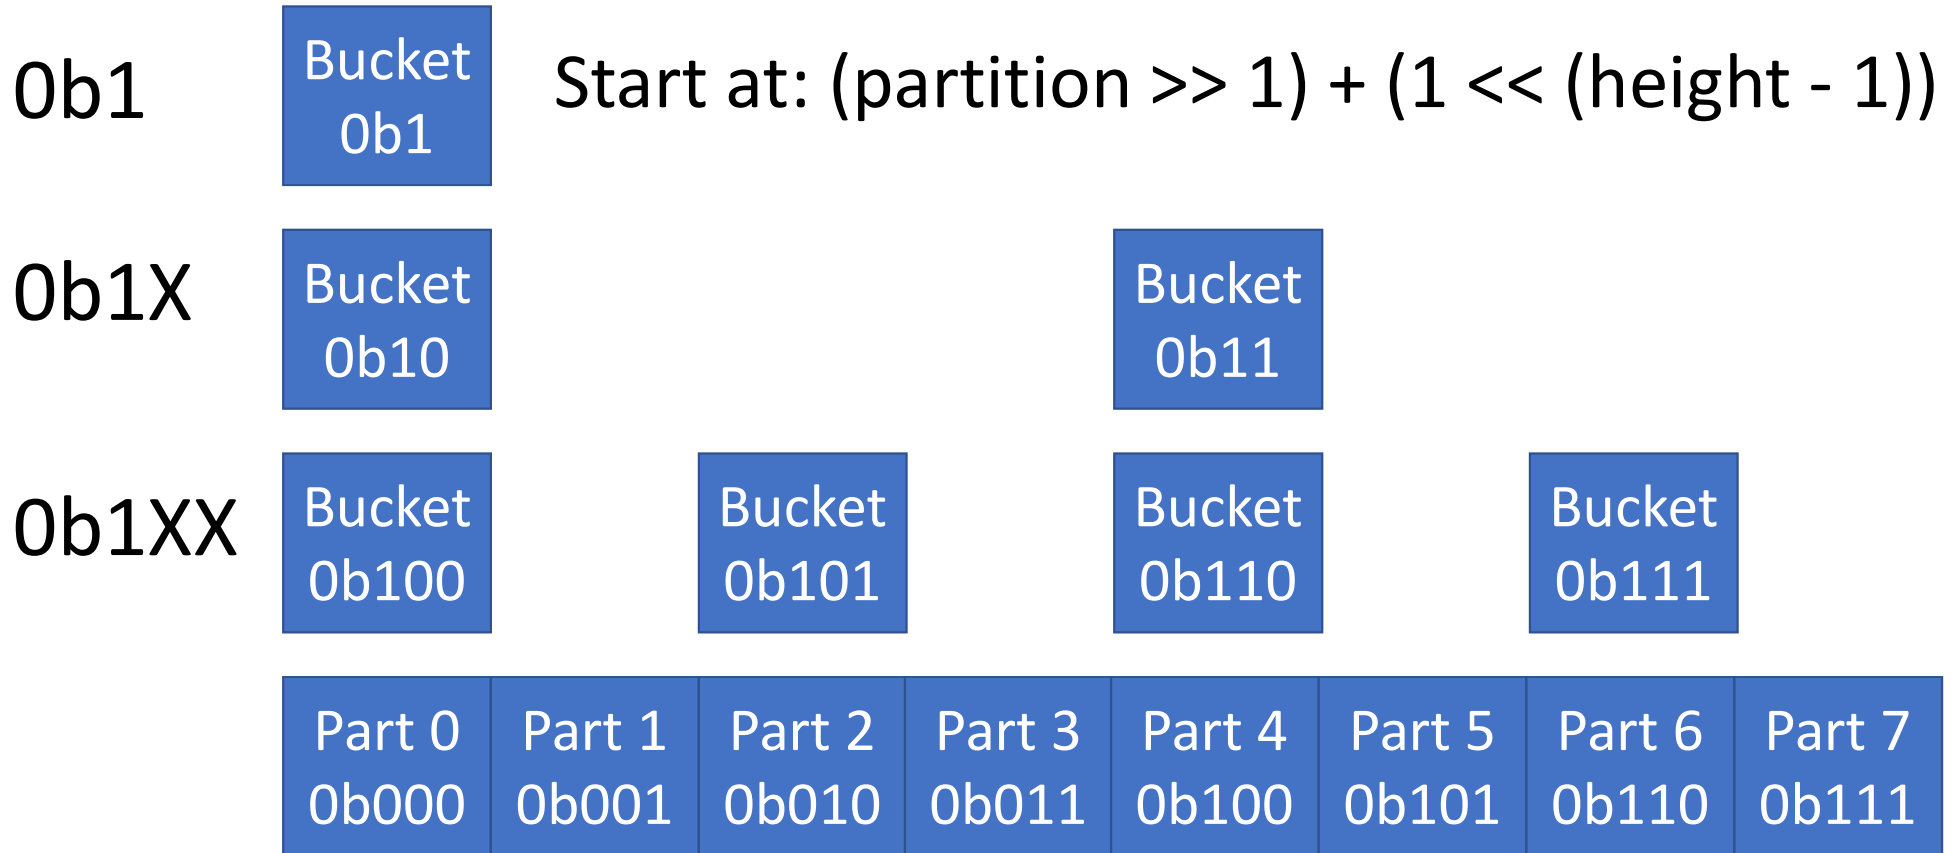
#partitions - 1 buckets



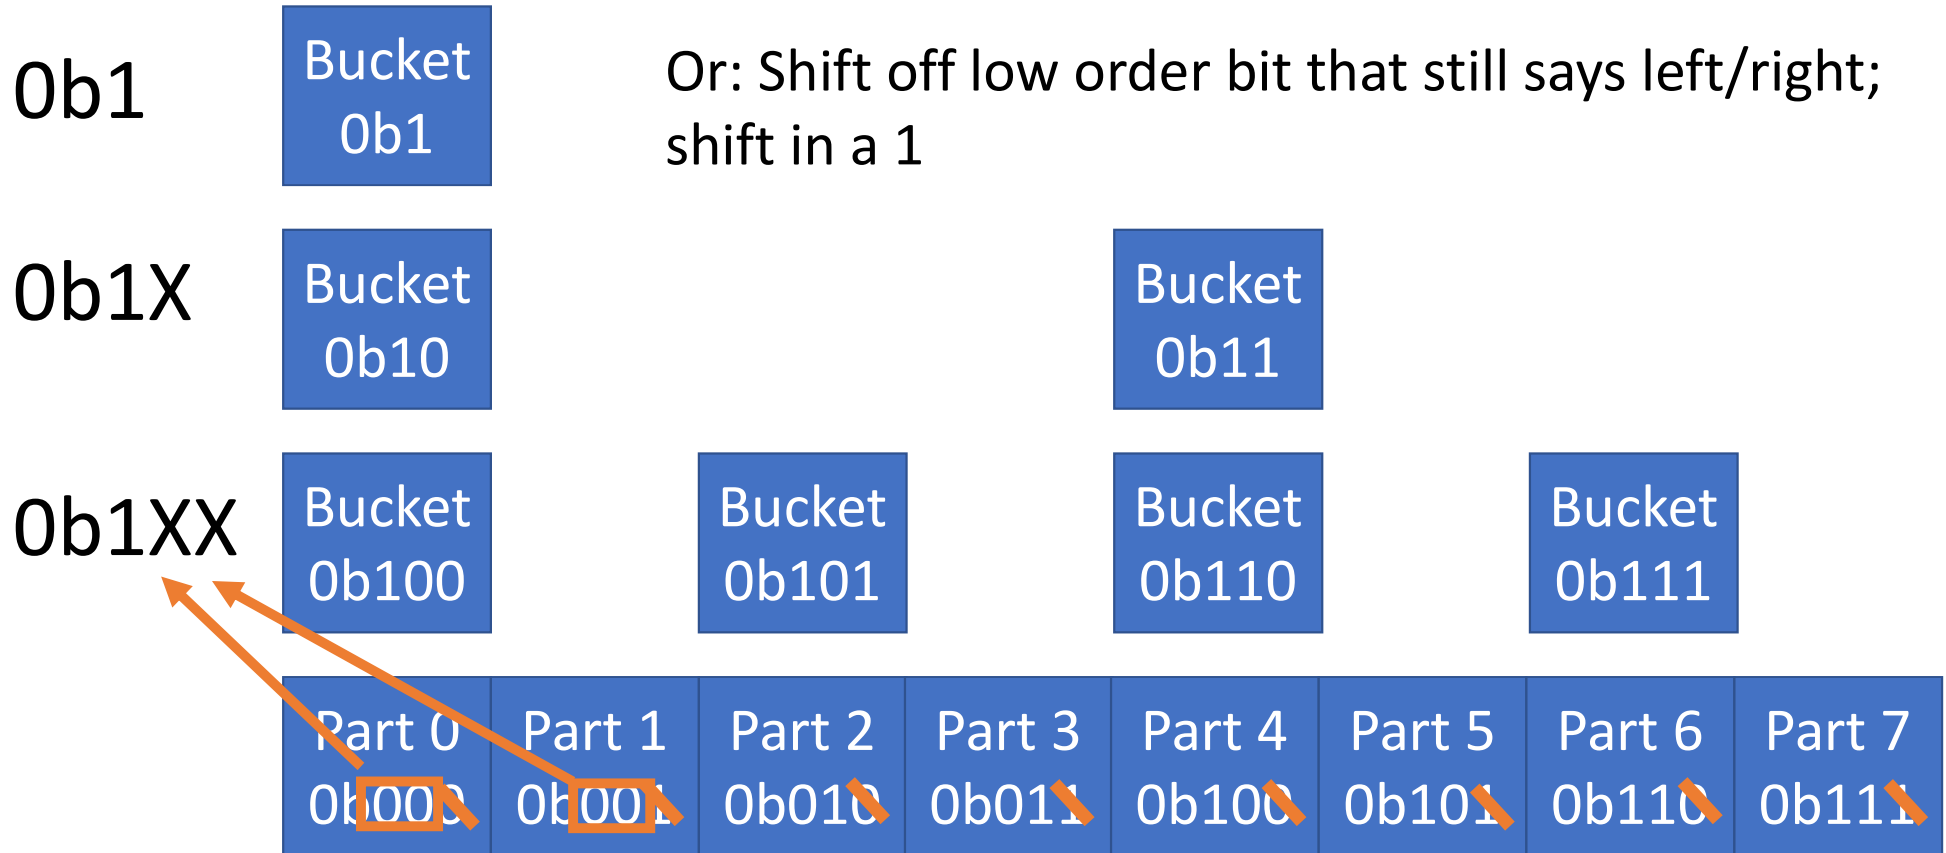
Bottom-up merge tree



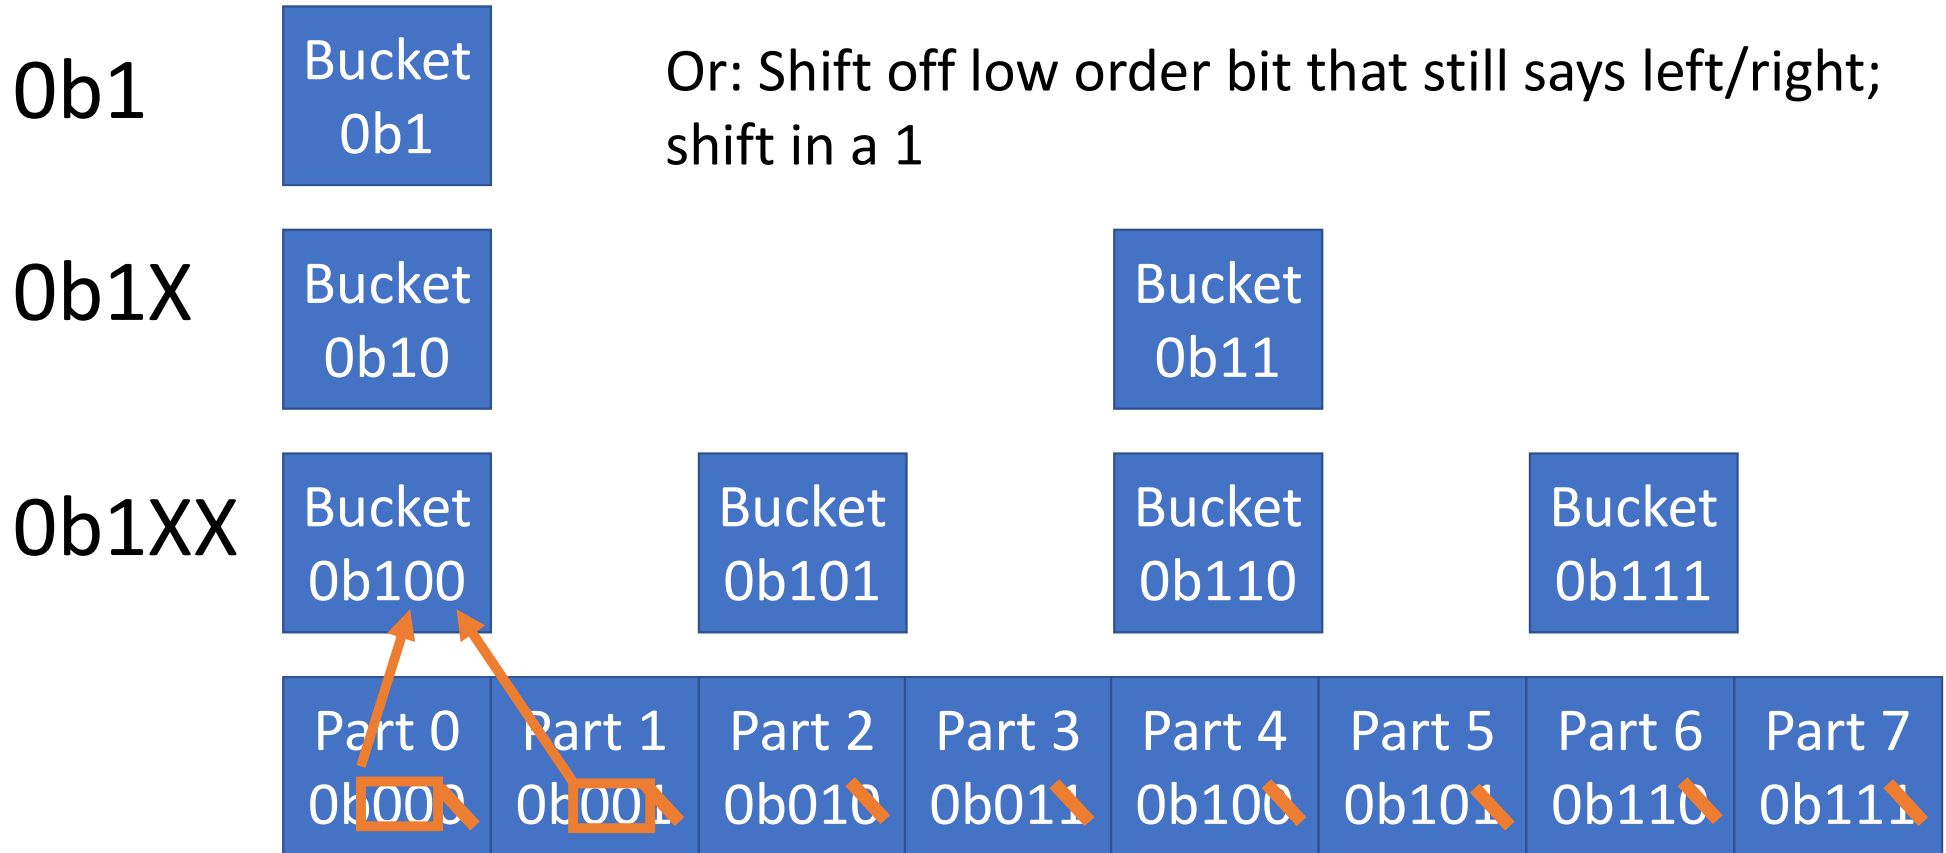
Bottom-up merge tree



Bottom-up merge tree




Bottom-up merge tree



stable_sort laptop results

- `.\stable_sort_release.exe 1000000000` - Parallel was 3.40229 times faster
- `.\stable_sort_debug.exe 1000000` - Parallel was 3.68103 times faster
- `.\stable_sort_release.exe 1000000` - Parallel was 2.96105 times faster
- `.\stable_sort_debug.exe 10000` - Parallel was 1.26545 times faster
- `.\stable_sort_release.exe 10000` - Parallel was 1.28868 times slower
- `.\stable_sort_debug.exe 100` - Parallel was 2.55325 times slower
- `.\stable_sort_release.exe 100` - Parallel was 10.2388 times slower

Take our survey <https://aka.ms/cppcon>



You can win an Xbox One S - Starter Bundle



9/27 10:30 – 12:00 // Breckenridge Hall

Thoughts on a More Powerful and Simpler C++ (5 of N), *Herb Sutter*

Other sessions

Monday, September 24th

- 14:00 – 15:00
How to Write Well-Behaved Value Wrappers
 - by Simon Brand
- 15:15 – 16:15
How C++ Debuggers Work
 - by Simon Brand

Tuesday, September 25th

- 14:00 – 15:00
What Could Possibly Go Wrong?: A Tale of Expectations and Exceptions
 - by Simon Brand and Phil Nash
- 15:15 – 15:45
Overloading: The Bane of All Higher-Order Functions
 - by Simon Brand

Wednesday, September 26th

- 12:30 – 13:30
C++ Community Building Birds of a Feather
 - with Stephan T. Lavavej and others
- 14:00 – 15:00
Latest and Greatest in the Visual Studio Family for C++ Developers 2018
 - by Marian Luparu and Steve Carroll
- 15:15 – 15:45
Don't Package Your Libraries, Write Packagable Libraries!
 - by Robert Schumacher

Wednesday, September 26th

- 15:15 – 15:45
What's new in Visual Studio Code for C++ Development
 - by Rong Lu
- 15:50 – 16:20
Value Semantics: Fast, Safe, and Correct by Default
 - by Nicole Mazzuca
- 16:45 – 17:45
Memory Latency Troubles You? Nano-coroutines to the Rescue! (Using Coroutines TS, of Course)
 - by Gor Nishanov
- 18:45 – 20:00
Cross-Platform C++ Development is Challenging – let Tools Help!
 - by Marc Goodner and Will Buik

Thursday, September 27th

- 9:00 – 10:00
Inside Visual C++'s Parallel Algorithms
 - by Billy O'Neal
- 15:15 – 15:45
ConcurrencyCheck – Static Analyzer for Concurrency Issues in Modern C++
 - by Anna Gringauze
- 16:45 – 17:45
Class Template Argument Deduction for Everyone
 - by Stephan T. Lavavej

Questions?

References for those looking at doing their own implementations:

A Parallel Algorithm for the Efficient Solution of a General Class of Recurrence Equations

Peter M. Kogge and Harold S. Stone

Single-pass Parallel Prefix Scan with Decoupled Look-back

Duane Merrill and Michael Garland

Thread Scheduling for Multiprogrammed Multiprocessors

Nimar S. Arora, Robert D. Blumofe, and C. Greg Plaxton

Dynamic Circular Work-Stealing Deque

David Chase and Yossi Lev

CppCon 2015 "Work Stealing" <https://youtu.be/iLHNF7SgVN4>

Pablo Halpern

Inside Windows 8: Pedro Teixeira - Thread pools [https://channel9.msdn.com/Shows/Going+Deep/Inside-Windows-8-](https://channel9.msdn.com/Shows/Going+Deep/Inside-Windows-8-Pedro-Teixeira-Thread-pool)

[Pedro-Teixeira-Thread-pool](https://channel9.msdn.com/Shows/Going+Deep/Inside-Windows-8-Pedro-Teixeira-Thread-pool)

Pedro Teixeira